

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.arima_model import ARMA
from sklearn.svm import LinearSVC
```

```
In [2]: df = pd.read_csv('fullset2.csv', index_col='timestamp')
print(df.head())
n_row, n_col = df.shape
print(f'There are {n_row} rows and {n_col} columns')
```

	amount	price	coin
timestamp			
2018-08-15 18:49:25	0.233800	529.41	1
2018-08-15 18:50:32	0.982765	529.45	1
2018-08-15 18:50:33	0.017235	529.45	1
2018-08-15 18:50:33	0.017235	529.46	1
2018-08-15 18:50:34	0.965530	529.45	1

Let's add a new column calculating the amount of cryptocurrency purchased with the daily price and call the new column Value.

```
In [3]: df['value'] = df.amount * df.price
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1314830 entries, 2018-08-15 18:49:25 to 2018-09-22 10:07:01
Data columns (total 4 columns):
amount      1314830 non-null float64
price       1314830 non-null float64
coin        1314830 non-null int64
value       1314830 non-null float64
dtypes: float64(3), int64(1)
memory usage: 50.2+ MB
```

If we want to change the coin labels from 1,2,3,4,5 to their respective names, we can use the .apply() method and lambda functions over the coin column. But First we'll need to change the column type for string manipulation.

```
In [4]: # Lamda function using replace
df['coin'] = df.coin.astype(str)
df['coin'] =df.coin.apply(lambda x: x.replace('1', 'Bitcoin Cash'))
df['coin'] =df.coin.apply(lambda x: x.replace('2', 'Bitcoin'))
df['coin'] =df.coin.apply(lambda x: x.replace('3', 'Ethereum'))
df['coin'] =df.coin.apply(lambda x: x.replace('4', 'Litecoin'))
df['coin'] =df.coin.apply(lambda x: x.replace('5', 'Ripple'))
assert df.coin.dtypes == np.object
print(df.tail())
```

	amount	price	coin	value
timestamp				
2018-09-22 10:06:34	95.389900	0.58073	Ripple	55.395777
2018-09-22 10:06:34	253.735000	0.58072	Ripple	147.348989
2018-09-22 10:06:35	2150.875100	0.58000	Ripple	1247.507558
2018-09-22 10:06:50	5133.932400	0.58201	Ripple	2987.999996
2018-09-22 10:07:01	222.319377	0.58326	Ripple	129.670000

```
In [5]: print(df.describe())
print(df['coin'].value_counts(dropna=False))
print(df['price'].describe())
print(df['price'].max)
print(df['coin'].min)
```

	amount	price	value
count	1.314830e+06	1.314830e+06	1.314830e+06
mean	5.372889e+02	3.130907e+03	1.588603e+03
std	4.234103e+03	3.215822e+03	6.956488e+03
min	1.000000e-08	2.530000e-01	2.557900e-09
25%	3.194035e-02	6.047000e+01	2.411842e+01
50%	3.300000e-01	4.715600e+02	2.030175e+02
75%	6.495622e+00	6.462340e+03	1.018026e+03
max	8.995980e+05	7.411850e+03	1.764929e+06

Bitcoin 607607

Ethereum 290479

Ripple 270182

Litecoin 74137

Bitcoin Cash 72425

Name: coin, dtype: int64

count 1.314830e+06

mean 3.130907e+03

std 3.215822e+03

min 2.530000e-01

25% 6.047000e+01

50% 4.715600e+02

75% 6.462340e+03

max 7.411850e+03

Name: price, dtype: float64

<bound method Series.max of timestamp

2018-08-15 18:49:25 529.41000

2018-08-15 18:50:32 529.45000

2018-08-15 18:50:33 529.45000

2018-08-15 18:50:33 529.46000

2018-08-15 18:50:34 529.45000

2018-08-15 18:50:34 529.46000

2018-08-15 18:55:21 528.38000

2018-08-15 18:55:29 528.38000

2018-08-15 18:55:37 527.99000

2018-08-15 18:55:37 527.94000

2018-08-15 18:56:24 528.92000

2018-08-15 19:00:00 529.01000

2018-08-15 19:00:01 529.76000

2018-08-15 19:00:57 529.77000

2018-08-15 19:07:31 528.00000

2018-08-15 19:07:34 527.52000

2018-08-15 19:08:48 526.33000

2018-08-15 19:08:50 525.69000

2018-08-15 19:08:52 525.69000

2018-08-15 19:08:52 525.69000

2018-08-15 19:08:53 525.59000

2018-08-15 19:08:53 525.39000

2018-08-15 19:08:53 525.17000

2018-08-15 19:08:54 526.32000

2018-08-15 19:09:09 525.06000

2018-08-15 19:09:52 525.04000

2018-08-15 19:10:11 525.70000

2018-08-15 19:12:49 525.74000

2018-08-15 19:13:41 525.72000

2018-08-15 19:14:11 525.59000

...

2018-09-22 10:03:51 0.58000

```

2018-09-22 10:03:51    0.58001
2018-09-22 10:03:51    0.58010
2018-09-22 10:03:52    0.58010
2018-09-22 10:03:52    0.58040
2018-09-22 10:03:52    0.58040
2018-09-22 10:03:53    0.58050
2018-09-22 10:03:53    0.58100
2018-09-22 10:03:53    0.58105
2018-09-22 10:03:53    0.58105
2018-09-22 10:03:54    0.58110
2018-09-22 10:03:58    0.58170
2018-09-22 10:03:58    0.58168
2018-09-22 10:04:01    0.58249
2018-09-22 10:04:01    0.58249
2018-09-22 10:04:01    0.58251
2018-09-22 10:04:19    0.58249
2018-09-22 10:05:04    0.57938
2018-09-22 10:05:04    0.57938
2018-09-22 10:05:04    0.57938
2018-09-22 10:05:26    0.57940
2018-09-22 10:05:43    0.58140
2018-09-22 10:05:44    0.57781
2018-09-22 10:05:57    0.58220
2018-09-22 10:06:00    0.58201
2018-09-22 10:06:34    0.58073
2018-09-22 10:06:34    0.58072
2018-09-22 10:06:35    0.58000
2018-09-22 10:06:50    0.58201
2018-09-22 10:07:01    0.58326

```

Name: price, Length: 1314830, dtype: float64>

<bound method Series.min of timestamp

```

2018-08-15 18:49:25    Bitcoin Cash
2018-08-15 18:50:32    Bitcoin Cash
2018-08-15 18:50:33    Bitcoin Cash
2018-08-15 18:50:33    Bitcoin Cash
2018-08-15 18:50:34    Bitcoin Cash
2018-08-15 18:50:34    Bitcoin Cash
2018-08-15 18:55:21    Bitcoin Cash
2018-08-15 18:55:29    Bitcoin Cash
2018-08-15 18:55:37    Bitcoin Cash
2018-08-15 18:55:37    Bitcoin Cash
2018-08-15 18:56:24    Bitcoin Cash
2018-08-15 19:00:00    Bitcoin Cash
2018-08-15 19:00:01    Bitcoin Cash
2018-08-15 19:00:57    Bitcoin Cash
2018-08-15 19:07:31    Bitcoin Cash
2018-08-15 19:07:34    Bitcoin Cash
2018-08-15 19:08:48    Bitcoin Cash
2018-08-15 19:08:50    Bitcoin Cash
2018-08-15 19:08:52    Bitcoin Cash
2018-08-15 19:08:52    Bitcoin Cash
2018-08-15 19:08:53    Bitcoin Cash
2018-08-15 19:08:53    Bitcoin Cash
2018-08-15 19:08:53    Bitcoin Cash
2018-08-15 19:08:54    Bitcoin Cash
2018-08-15 19:09:09    Bitcoin Cash
2018-08-15 19:09:52    Bitcoin Cash

```

```

2018-08-15 19:10:11 Bitcoin Cash
2018-08-15 19:12:49 Bitcoin Cash
2018-08-15 19:13:41 Bitcoin Cash
2018-08-15 19:14:11 Bitcoin Cash
...
2018-09-22 10:03:51 Ripple
2018-09-22 10:03:51 Ripple
2018-09-22 10:03:51 Ripple
2018-09-22 10:03:52 Ripple
2018-09-22 10:03:52 Ripple
2018-09-22 10:03:52 Ripple
2018-09-22 10:03:53 Ripple
2018-09-22 10:03:53 Ripple
2018-09-22 10:03:53 Ripple
2018-09-22 10:03:53 Ripple
2018-09-22 10:03:54 Ripple
2018-09-22 10:03:58 Ripple
2018-09-22 10:03:58 Ripple
2018-09-22 10:04:01 Ripple
2018-09-22 10:04:01 Ripple
2018-09-22 10:04:01 Ripple
2018-09-22 10:04:19 Ripple
2018-09-22 10:05:04 Ripple
2018-09-22 10:05:04 Ripple
2018-09-22 10:05:04 Ripple
2018-09-22 10:05:26 Ripple
2018-09-22 10:05:43 Ripple
2018-09-22 10:05:44 Ripple
2018-09-22 10:05:57 Ripple
2018-09-22 10:06:00 Ripple
2018-09-22 10:06:34 Ripple
2018-09-22 10:06:34 Ripple
2018-09-22 10:06:35 Ripple
2018-09-22 10:06:50 Ripple
2018-09-22 10:07:01 Ripple
Name: coin, Length: 1314830, dtype: object>

```

Here, we can see Bitcoin Cash had the highest price value at 529.46 while Ripple was lowest at 0.58, coinciding with the head and tail ends of our dataset. Looking at some of the min/max values, we might want to investigate further for potential outliers.

```
In [6]: print(df.isnull().sum())  
        print(df.notnull())
```

```

amount    0
price     0
coin      0
value     0
dtype: int64

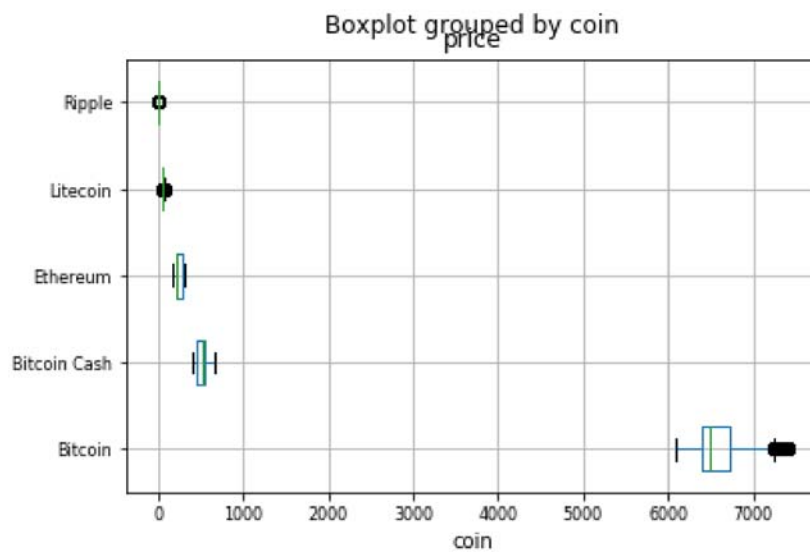
```

	amount	price	coin	value
timestamp				
2018-08-15 18:49:25	True	True	True	True
2018-08-15 18:50:32	True	True	True	True
2018-08-15 18:50:33	True	True	True	True
2018-08-15 18:50:33	True	True	True	True
2018-08-15 18:50:34	True	True	True	True
2018-08-15 18:50:34	True	True	True	True
2018-08-15 18:55:21	True	True	True	True
2018-08-15 18:55:29	True	True	True	True
2018-08-15 18:55:37	True	True	True	True
2018-08-15 18:55:37	True	True	True	True
2018-08-15 18:56:24	True	True	True	True
2018-08-15 19:00:00	True	True	True	True
2018-08-15 19:00:01	True	True	True	True
2018-08-15 19:00:57	True	True	True	True
2018-08-15 19:07:31	True	True	True	True
2018-08-15 19:07:34	True	True	True	True
2018-08-15 19:08:48	True	True	True	True
2018-08-15 19:08:50	True	True	True	True
2018-08-15 19:08:52	True	True	True	True
2018-08-15 19:08:52	True	True	True	True
2018-08-15 19:08:53	True	True	True	True
2018-08-15 19:08:53	True	True	True	True
2018-08-15 19:08:53	True	True	True	True
2018-08-15 19:08:54	True	True	True	True
2018-08-15 19:09:09	True	True	True	True
2018-08-15 19:09:52	True	True	True	True
2018-08-15 19:10:11	True	True	True	True
2018-08-15 19:12:49	True	True	True	True
2018-08-15 19:13:41	True	True	True	True
2018-08-15 19:14:11	True	True	True	True
...
2018-09-22 10:03:51	True	True	True	True
2018-09-22 10:03:51	True	True	True	True
2018-09-22 10:03:51	True	True	True	True
2018-09-22 10:03:52	True	True	True	True
2018-09-22 10:03:52	True	True	True	True
2018-09-22 10:03:52	True	True	True	True
2018-09-22 10:03:53	True	True	True	True
2018-09-22 10:03:53	True	True	True	True
2018-09-22 10:03:53	True	True	True	True
2018-09-22 10:03:53	True	True	True	True
2018-09-22 10:03:54	True	True	True	True
2018-09-22 10:03:58	True	True	True	True
2018-09-22 10:03:58	True	True	True	True
2018-09-22 10:04:01	True	True	True	True
2018-09-22 10:04:01	True	True	True	True
2018-09-22 10:04:01	True	True	True	True
2018-09-22 10:04:19	True	True	True	True
2018-09-22 10:05:04	True	True	True	True
2018-09-22 10:05:04	True	True	True	True

2018-09-22 10:05:04	True	True	True	True
2018-09-22 10:05:26	True	True	True	True
2018-09-22 10:05:43	True	True	True	True
2018-09-22 10:05:44	True	True	True	True
2018-09-22 10:05:57	True	True	True	True
2018-09-22 10:06:00	True	True	True	True
2018-09-22 10:06:34	True	True	True	True
2018-09-22 10:06:34	True	True	True	True
2018-09-22 10:06:35	True	True	True	True
2018-09-22 10:06:50	True	True	True	True
2018-09-22 10:07:01	True	True	True	True


[1314830 rows x 4 columns]

```
In [7]: df.boxplot(column='price', by='coin', vert=False, fontsize=8)
plt.show()
```



From the boxplot, we can see some potential outliers for Bitcoin. However, given the volatile nature of Bitcoin (and cryptos in general), these outliers could actually be valid data points. Again, further investigation is needed to determine whether or not to drop any of those points.

```
In [8]: df.plot(kind='scatter', x='amount', y='value', rot=70)
plt.show()
```



With the cryptocurrencies listed in an untidy format, let's use `.pivot_table()` method to create a new column for each unique value ("price") in a specified column ("type of coin"). We will also need to use `.reset_index()` to get back the original DataFrame after being pivoted.

```
In [9]: df = df.pivot_table(index=['timestamp'], columns='coin', values='price', aggfunc='last')
df = df.reset_index()
print(df.head())
print(f'There are {n_row} rows and {n_col} columns')
```

	coin	timestamp	Bitcoin	Bitcoin Cash	Ethereum	Litecoin	Ripple
0	2018-08-15	18:45:56	NaN	NaN	297.56	NaN	NaN
1	2018-08-15	18:46:15	6548.0	NaN	NaN	NaN	NaN
2	2018-08-15	18:46:24	6548.0	NaN	NaN	NaN	NaN
3	2018-08-15	18:46:38	6547.3	NaN	NaN	NaN	NaN
4	2018-08-15	18:46:41	NaN	NaN	NaN	NaN	0.29083


```
In [10]: df.describe()
```

```
Out[10]:
```


coin	Bitcoin	Bitcoin Cash	Ethereum	Litecoin	Ripple
count	345822.000000	46255.000000	174389.000000	54694.000000	173701.000000
mean	6595.253606	515.545627	243.616883	57.463469	0.350847
std	302.570501	63.628698	40.608975	4.122115	0.095184
min	6094.470000	407.700000	167.000000	47.090000	0.253000
25%	6385.960000	460.370000	208.730000	54.800000	0.287010
50%	6480.310000	517.720000	229.360000	56.600000	0.325500
75%	6738.417500	553.990000	284.770000	59.760000	0.346830
max	7410.800000	660.070000	321.180000	69.400000	0.764400

We can now see that for all coins except Bitcoin Cash, the distributions are negatively skewed (left-skewed). Bitcoin Cash is approximately a normal distribution with very little skewness. Let's see a histogram to verify.


```
In [11]: # Due to the much higher/lower values of bitcoin and Ripple, we can plot it se
parately
# ax = df['Bitcoin'].plot(kind='hist', bins=60, fontsize=6, figsize=(6,4))
ax = df['Bitcoin Cash'].plot(kind='hist', bins=40, fontsize=6, figsize=(10,6))
ax = df['Ethereum'].plot(kind='hist', bins=40, fontsize=6, figsize=(10,6))
ax = df['Litecoin'].plot(kind='hist', bins=40, fontsize=6, figsize=(10,6))
ax.set_xlabel('Prices', fontsize=10)
ax.set_ylabel('Histogram of Cryptocurrencies', fontsize=10)
plt.legend(fontsize=10)
plt.show()
```



```
In [12]: ax = df['Bitcoin'].plot(kind='hist', bins=60, fontsize=6, figsize=(6,4))
```



```
In [13]: ax = df['Ripple'].plot(kind='hist', bins=60, fontsize=6, figsize=(6,4))
```



We decided to use the last non-NaN value instance to fill in missing values as it most accurately resembles the flow of time series data. Other ways of imputing include filling missing values using the `mean()` or `mode()`.

```
In [14]: print(df.isnull().sum())  
         print(df.notnull())
```

```

coin
timestamp      0
Bitcoin        311574
Bitcoin Cash   611141
Ethereum       483007
Litecoin       602702
Ripple         483695

```

```
dtype: int64
```

coin	timestamp	Bitcoin	Bitcoin Cash	Ethereum	Litecoin	Ripple
0	True	False	False	True	False	False
1	True	True	False	False	False	False
2	True	True	False	False	False	False
3	True	True	False	False	False	False
4	True	False	False	False	False	True
5	True	False	False	True	False	False
6	True	True	False	False	False	False
7	True	True	False	False	False	False
8	True	False	False	True	False	True
9	True	False	False	True	False	True
10	True	True	False	False	False	False
11	True	True	False	False	False	False
12	True	True	False	False	False	False
13	True	True	False	True	False	True
14	True	False	False	True	False	False
15	True	True	False	False	False	False
16	True	False	False	True	False	False
17	True	False	False	True	False	True
18	True	False	False	True	False	True
19	True	False	False	False	False	True
20	True	False	False	True	False	True
21	True	True	False	False	False	False
22	True	True	False	True	False	False
23	True	True	False	False	False	False
24	True	True	False	False	False	False
25	True	True	False	False	False	False
26	True	True	False	False	False	False
27	True	False	False	True	False	False
28	True	True	False	False	False	False
29	True	True	False	False	False	False
...
657366	True	False	True	False	True	False
657367	True	False	False	True	False	False
657368	True	False	False	False	False	True
657369	True	False	False	False	False	True
657370	True	True	False	False	False	False
657371	True	False	False	True	False	False
657372	True	True	False	False	False	False
657373	True	False	False	False	False	True
657374	True	True	False	False	False	False
657375	True	True	True	False	False	False
657376	True	False	False	False	False	True
657377	True	False	False	True	False	False
657378	True	False	False	True	False	False
657379	True	False	False	True	False	False
657380	True	True	False	False	False	False
657381	True	False	False	False	True	False
657382	True	True	False	False	False	False

657383	True	True	False	True	False	False
657384	True	True	False	True	False	False
657385	True	False	True	False	True	False
657386	True	True	False	False	False	False
657387	True	False	False	True	False	False
657388	True	False	False	False	False	True
657389	True	False	False	False	False	True
657390	True	False	False	False	False	True
657391	True	False	False	True	False	False
657392	True	False	False	False	False	True
657393	True	True	False	False	False	False
657394	True	True	False	True	True	False
657395	True	False	True	False	False	False

[657396 rows x 6 columns]

In [15]: *#fill missing values with next instance since some initial values are missing*

```
df['Bitcoin Cash'] = df['Bitcoin Cash'].fillna(method='bfill')
df['Bitcoin'] = df['Bitcoin'].fillna(method='bfill')
df['Ethereum'] = df['Ethereum'].fillna(method='bfill')
df['Litecoin'] = df['Litecoin'].fillna(method='bfill')
df['Ripple'] = df['Ripple'].fillna(method='bfill')
```

In [16]: df.head(15)

Out[16]:

	coin	timestamp	Bitcoin	Bitcoin Cash	Ethereum	Litecoin	Ripple
0	2018-08-15 18:45:56	6548.00	529.41	297.56	58.04	0.29083	
1	2018-08-15 18:46:15	6548.00	529.41	297.81	58.04	0.29083	
2	2018-08-15 18:46:24	6548.00	529.41	297.81	58.04	0.29083	
3	2018-08-15 18:46:38	6547.30	529.41	297.81	58.04	0.29083	
4	2018-08-15 18:46:41	6547.30	529.41	297.81	58.04	0.29083	
5	2018-08-15 18:46:42	6547.30	529.41	297.81	58.04	0.29083	
6	2018-08-15 18:46:44	6547.30	529.41	297.81	58.04	0.29083	
7	2018-08-15 18:46:45	6550.96	529.41	297.81	58.04	0.29083	
8	2018-08-15 18:46:51	6547.42	529.41	297.81	58.04	0.29083	
9	2018-08-15 18:46:58	6547.42	529.41	297.81	58.04	0.29083	
10	2018-08-15 18:46:59	6547.42	529.41	297.81	58.04	0.29080	
11	2018-08-15 18:47:00	6555.00	529.41	297.81	58.04	0.29080	
12	2018-08-15 18:47:04	6547.18	529.41	297.81	58.04	0.29080	
13	2018-08-15 18:47:05	6554.57	529.41	297.81	58.04	0.29080	
14	2018-08-15 18:47:06	6552.92	529.41	297.81	58.04	0.29084	

```
In [17]: print(df.dtypes)
```

```
coin
timestamp    object
Bitcoin      float64
Bitcoin Cash float64
Ethereum     float64
Litecoin     float64
Ripple       float64
dtype: object
```

Before delving into any further analysis, let's modify the 'timestamp' column to reflect datetime64 instead of it's current object type. This will allow for fascilitating time series analysis.

```
In [18]: df['timestamp'] = pd.to_datetime(df['timestamp'])

# Set the timestamp column as the dataframe index
df = df.set_index(['timestamp'])
```



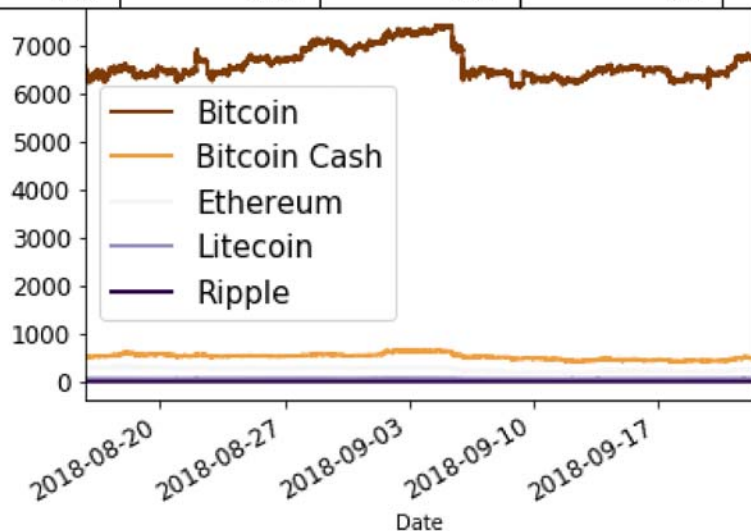
```
In [19]: ax = df.plot(colormap='PuOr',linewidth=2,fontsize=12)

# Additional customizations
ax.set_xlabel('Date')
ax.legend(fontsize=15)

df_summary = df.describe()
# Add summary statistics to plot
ax.table(cellText=df_summary.values,
         colWidths=[0.3]*len(df.columns),
         rowLabels=df_summary.index,
         colLabels=df_summary.columns,
         fontsize=15,
         loc='top')

# Show plot - plotted separately a couple boxes down to better discern plots o
f each coin
plt.show()
```


	Bitcoin	Bitcoin Cash	Ethereum	Litecoin	Ripple
count	657395.0	657396.0	657395.0	657395.0	657393.0
mean	6584.876537903439	510.58725798523926	249.45852323185093	57.25940223152117	0.32841412270894266
std	298.47820223356194	56.21462392469781	39.85892181626725	3.858675011369461	0.0694407334833036
min	6094.47	407.7	167.0	47.09	0.253
25%	6380.68	464.48	211.85	54.79	0.28316
50%	6477.7	518.88	262.48	56.58	0.32267
75%	6727.17	545.47	286.0	59.04	0.33815
max	7410.8	660.07	321.18	69.4	0.7644



```
In [20]: ax = df.plot.area(fontsize=12)

# Additional customizations
ax.set_xlabel('Date')
ax.legend(fontsize=15)

# Show plot
plt.show()
```




We are going to focus and mark the time period (denoted by green dashlines) when Bitcoin prices drop into a slight valley. Prior to the drop, Bitcoin was very steadily increasing at a positive rate.

```
In [21]: ax = df.plot(figsize=(9,5), linewidth=1, fontsize=6, colormap='Spectral')
ax.set_xlabel('Time Stamp')
ax.set_ylabel('Price')

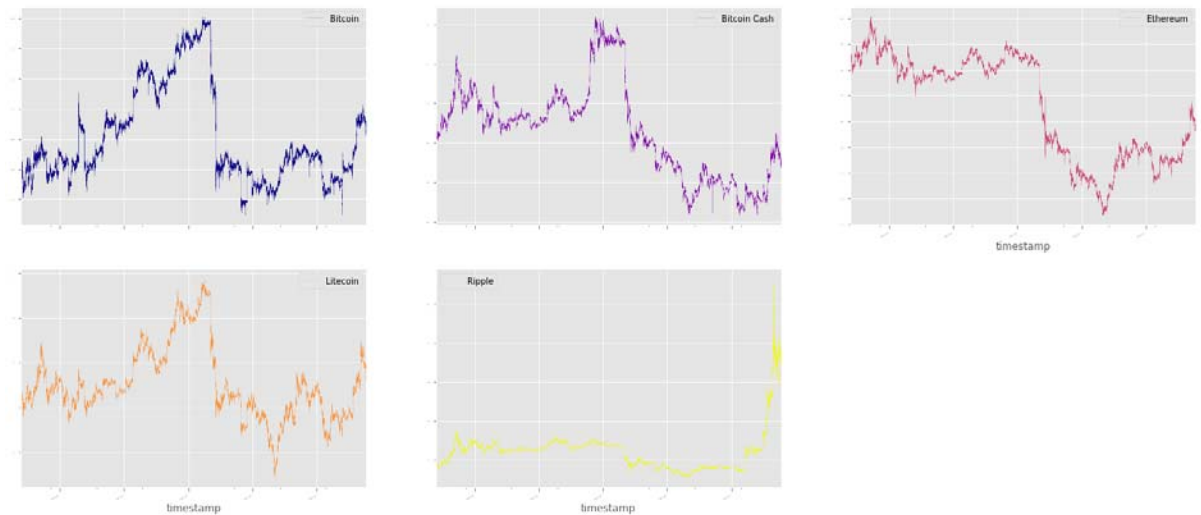
#highlight the time periods where prices dipped denoted by dashed green vertical lines
ax.axvline('2018-09-04 23:5:56', color='green', linestyle='--')
ax.axvline('2018-09-21', color='green', linestyle='--')

plt.style.use('ggplot')
ax.set_title('Crypto Prices')
ax.legend(loc='center left', bbox_to_anchor=(1.0, 0.5))
plt.show()
```



Right now, it is hard to discern the data trends for any other currency besides Bitcoin. So, let's create subplots for each coin to improve clarity and provide more context.

```
In [22]: df.plot(subplots=True,  
                 layout= (3,3),  
                 sharex=True,  
                 sharey=False,  
                 colormap='plasma',  
                 fontsize=2,  
                 legend=True,  
                 linewidth=0.2,  
                 figsize=(25,18))  
  
plt.show()
```



We can observe that for most of the coins, around 09-04-2018 (~middle of x-axis) the prices all similarly dropped. Appears there is at least a direct negative correlation as all currencies were affected by a negative drop in value.

Now that we can view the visualizations for each individual coin, we want to determine if and how the coins are correlated with one another. Here, we'll be computing the correlation coefficients using pearson and spearman, depending on whether the relationships are thought to be linear or not.


```
In [23]: corr_bitlite = df[['Bitcoin', 'Litecoin']].corr(method='spearman')
print(df[['Bitcoin', 'Litecoin']].corr(method='spearman'))
print("The correlation between Bitcoin and Bitcoin Cash is 0.909616", '\n')

corr_df = df.corr(method='spearman')

#create heatmap of correlation matrix
sns.heatmap(corr_df,
            annot=True,
            linewidths=0.4,
            annot_kws={"size": 10})

plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.show()
```

```
coin      Bitcoin  Litecoin
coin
Bitcoin    1.000000  0.909616
Litecoin   0.909616  1.000000
('The correlation between Bitcoin and Bitcoin Cash is 0.909616', '\n')
```




Here we see using the Spearman method, though not necessarily ordinal values, we can use this method to rank each coin and correlation coefficients, which can then be used to summarise the monotonic (entirely nonincreasing or nondecreasing) strength and direction of the relationships. Later, we will use the Pearson correlation method to determine the linear strength and direction between the variables. Comparing the Pearson graph below with the produced Spearman Graph, we can denote anytime correlation values for (S)pearman > (P)earson (Litecoin/Ethereum; Bitcoin Cash/Ripple), we have a correlation that is monotonic but not linear. With Pearson correlation higher, we can discern the linear correlation is larger than rank. This may be due to influential observations in the distribution tails having a large influence relative to ranked values. If linearity holds in our dataset, we can associate Pearson correlation to be of stronger measure.

Let's now visualize the correlation matrix using the pearson method.

```
In [24]: corr_df = df.corr(method='pearson')

#create heatmap of correlation matrix
sns.heatmap(corr_df,
            annot=True,
            linewidths=0.4,
            annot_kws={"size": 10})

plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.show()
```




Let's also visualize a clustermap using the pearson method.

```
In [25]: corr_df = df.corr(method='pearson')

# Customize the heatmap of the corr_meat correlation matrix and rotate the x-axis labels
fig = sns.clustermap(corr_df,
                     row_cluster=True,
                     col_cluster=True,
                     figsize=(10, 10))

plt.setp(fig.ax_heatmap.xaxis.get_majorticklabels(), rotation=90)
plt.setp(fig.ax_heatmap.yaxis.get_majorticklabels(), rotation=0)
plt.show()
```



Again, you can see a stronger correlations between Bitcoin/Litecoin.

Next, let's introduce some Time series decomposition as it can be a great way to reveal the time series structure. Let's obtain the seasonal decomposition and visualize the components

```
In [26]: # Couple of NaN values remain at the tail of df after backfilling values. Since seasonal_decompose doesn't handle missing values, we will drop all NaN.  
df = df.dropna()
```

```
In [27]: # import seasonal_decompose() function from statsmodels library  
import statsmodels.api as sm  
  
# Initialize dictionary  
coin_decomp = {}  
  
# Get the names of each time series in the DataFrame  
coin_names = df.columns  
  
# Run time series decomposition on each time series of the DataFrame  
for ts in coin_names:  
    ts_decomposition = sm.tsa.seasonal_decompose(df[ts].values, freq=10)  
    coin_decomp[ts] = ts_decomposition
```


```
In [28]: # Extract the seasonal values for the decomposition of each time series  
coin_seasonal = {}  
for ts in coin_names:  
    coin_seasonal[ts] = coin_decomp[ts].seasonal  
  
# Create a DataFrame from the jobs_seasonal dictionary  
seasonality_df = pd.DataFrame.from_dict(coin_seasonal)  
  
# # Remove the Label for the index  
# seasonality_df.index.name = None  
  
# plot of the seasonality_df DataFrame  
# seasonality_df.plot(subplots=True,  
#                      layout=(4,4),  
#                      sharey=False,  
#                      fontsize=2,  
#                      linewidth=0.3,  
#                      legend=False,  
#                      figsize=(20,16))  
  
# # Show plot  
# plt.show()
```



```
In [29]: # Get correlation matrix of the seasonality_df DataFrame
seasonality_corr = seasonality_df.corr(method='spearman')

# Customize the clustermap of the seasonality_corr correlation matrix
fig = sns.clustermap(seasonality_corr, annot=True, annot_kws={"size": 16}, line
widths=.4, figsize=(12, 8))
plt.setp(fig.ax_heatmap.yaxis.get_majorticklabels(), rotation=0)
plt.setp(fig.ax_heatmap.xaxis.get_majorticklabels(), rotation=90)
plt.title('Seasonlity Trends Correlation Matrix')
plt.show()
```


Seasonlity Trends Correlation Matrix



From the seasonality trends matrix above, we can say Ripple is negatively correlated with Bitcoin Cash (-.71), Ripple has some positive correlation with Ethereum (0.56), and there does not seem to be much correlation between Bitcoin with Litecoin (-0.15). * Due to the 1-year timeframe of our dataset, the seasonlity trends would certainly make for stronger analysis on a larger timeframe.

Taking a closer look at just the Bitcoin trend, let's revisit and highlight the steep drop occuring at 2018-09-04.

```
In [30]: ax = df['Bitcoin'].plot(color='red', figsize=(9,5), linewidth=2, fontsize=6)
ax.set_xlabel('Time Stamp')
ax.set_ylabel('Price')
plt.style.use('ggplot')
ax.set_title('Crypto Prices')
ax.axvline('2018-09-04 14:45:56', color='green', linestyle='--')
ax.axhline(7410.8, color='green',linestyle='--')
plt.show()
```



Looking at the summary statistics and the graph above, we can see the max value for Bitcoin was 7410.80 occurring on 2018-09-04, before the price dropped a low of 6094.47. Also, below we've calculated the 5% and 95% values as 6250.62 and 7237.37 for Bitcoin.

```
In [31]: print(df.describe())
```

coin	Bitcoin	Bitcoin Cash	Ethereum	Litecoin \
count	657393.000000	657393.000000	657393.000000	657393.000000
mean	6584.876185	510.587407	249.458560	57.259398
std	298.478588	56.214709	39.858977	3.858680
min	6094.470000	407.700000	167.000000	47.090000
25%	6380.680000	464.480000	211.850000	54.790000
50%	6477.700000	518.880000	262.480000	56.580000
75%	6727.170000	545.470000	286.000000	59.040000
max	7410.800000	660.070000	321.180000	69.400000

coin	Ripple
count	657393.000000
mean	0.328414
std	0.069441
min	0.253000
25%	0.283160
50%	0.322670
75%	0.338150
max	0.764400

```
In [32]: # remove timestamp columns as it is not type int
filter_df = df.loc[:,df.columns != 'timestamp']

#Now let's compute the percentiles
low = 0.05
high = 0.95
quant_df = filter_df.quantile([low,high])
print(quant_df)

# filter values based on printed limits
filter_df = filter_df.apply(lambda x: x[(x>=quant_df.loc[low,x.name]) &
                                     (x <= quant_df.loc[high,x.name])], axis=0)


#drop any Nan values
filter_df.dropna(inplace=True)
```

coin	Bitcoin	Bitcoin Cash	Ethereum	Litecoin	Ripple
0.05	6250.620	426.06	189.35	52.07	0.26610
0.95	7237.368	626.89	299.35	65.51	0.51011

Again, let's center the boxplot from earlier to just Bitcoin to zoom in on whiskers and outlying data.


```
In [33]: fig = plt.figure(figsize=(8,6))
ax = fig.gca()
df.boxplot(column='Bitcoin', ax=ax)
# frame['ArrDelay'].plot.box(ax=ax) # Alternative
ax.set_title('Bitcoin Boxplot')
ax.set_xlabel('')
ax.set_ylabel('Prices')
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 657393 entries, 2018-08-15 18:45:56 to 2018-09-22 10:07:01
Data columns (total 5 columns):
Bitcoin          657393 non-null float64
Bitcoin Cash     657393 non-null float64
Ethereum         657393 non-null float64
Litecoin         657393 non-null float64
Ripple           657393 non-null float64
dtypes: float64(5)
memory usage: 30.1 MB
None
```




Again, to better visualize just the Bitcoin data, let's look at a density plot.

```
In [34]: ax = df['Bitcoin'].plot(kind='density', linewidth=4, fontsize=6, figsize=(10,6))
ax.set_xlabel('Prices', fontsize=10)
ax.set_ylabel('Density plot for Bitcoin', fontsize=10)
plt.show()
```



```
In [35]: import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')
from statsmodels.graphics import tsaplots

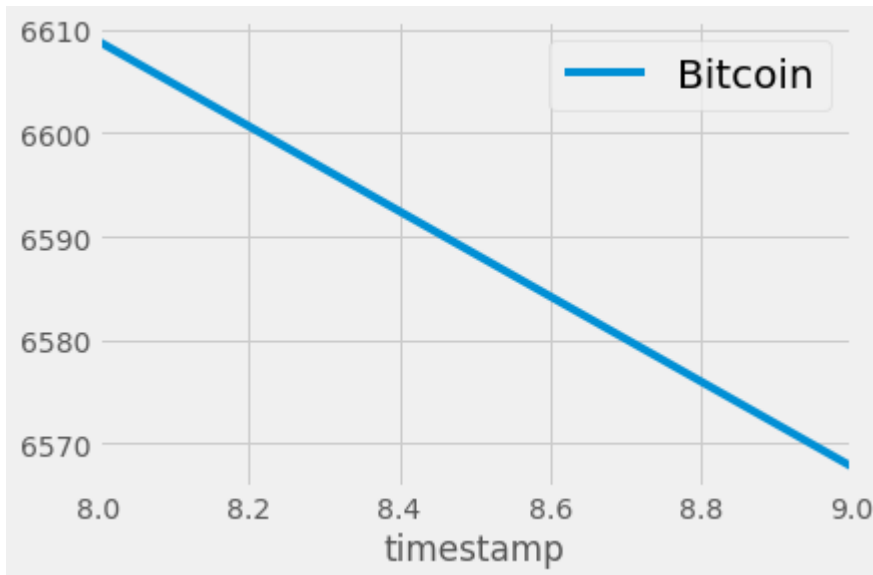
fig = tsaplots.plot_pacf(df['Bitcoin'], lags=30)
plt.show()
```



```
In [36]: index_month = df.index.month

# Compute the mean number of Bitcoin price from August to Sept
mean_price_by_month = df['Bitcoin'].groupby(index_month).mean()

mean_price_by_month.plot()
plt.legend(fontsize=20)
plt.show()
```




With autocorrelation values close to 0, we can conclude values between consecutive observations are not correlated with one another.

What if we wanted to be able to predict the future price of the coins based on our data and past trends? We would first need to split our data into train-test splits so we can test the quality of our model fit.

```
In [37]: coin_train = df.loc[:'2018-08']
coin_test = df.loc['2018-09':]
fig, ax = plt.subplots()

#Plot train and test sets
coin_train.plot(ax=ax)
coin_test.plot(ax=ax)
plt.legend(fontsize=8)
plt.show()
```



In order to statistically test whether the null hypothesis is our time series data is non-stationary due to trend, we can implement the augmented Dicky-Fuller test (this can be imported in Python as `adfuller`). If we can decide the data is non-stationary, then we will have to transform it into a stationary set prior to making our predictions. This can often be done by transforming the data by taking the difference, log, square root, or proportional change. We will certainly look to find the simplest yet effective implementation. This time, let's switch over to Litecoin column of the dataset to determine if it's stationary (p-value < 0.05 significance). If we wanted a p-value of 0.05 or below, the test statistic (ADF Statistic) needs to be below the 5% (or -2.8615 below) critical value of the test statistic (Reference: <https://machinelearningmastery.com/time-series-data-stationary-python/> (<https://machinelearningmastery.com/time-series-data-stationary-python/>))

```
In [38]: # Import augmented dicky-fuller test function
from statsmodels.tsa.stattools import adfuller

# Run test
result = adfuller(df['Litecoin'],2) #maxlag

# Print p-value
print('p-value:', result[1])

# Print critical values
print('critical values', result[4])

# Print the test statistic
print('ADF Statistic:', result[0])

('p-value:', 0.1555345950682367)
('critical values', {'5%': -2.8615443966388914, '1%': -3.4303599474064623, '10%': -2.5667723401698734})
('ADF Statistic:', -2.3526970834847893)
```

Next, we'll take the difference from each value in our time series by subtracting the previous value using the `.diff()` method. The results below shows the impact on our p-value as it drops to 0.0 or below 0.05 and the ADF statistic becomes more negative, stabilizing our data to a more stationary stance for potential future modeling (Again, using Dicky Fuller Statistic).


```

In [39]: # Calculate the difference of the time series
df_stationary = df.diff().dropna()

# Run ADF test on the differenced time series
result = adfuller(df_stationary['Litecoin'],2)

# Plot the differenced time series
fig, ax = plt.subplots()
df_stationary.plot(ax=ax)
plt.legend(fontsize=8)
plt.show()


# Print the test statistic and the p-value
print('ADF Statistic:', result[0])

# Print p-value
print('p-value:', result[1])

# Print critical values
print('critical values', result[4])

# Print the test statistic and the p-value
print('ADF Statistic:', result[0])

```



```

('ADF Statistic:', -496.54466639688314)
('p-value:', 0.0)
('critical values', {'5%': -2.8615443966455794, '1%': -3.4303599474215942, '10%': -2.566772340173433})
('ADF Statistic:', -496.54466639688314)

```

```
In [40]: ## Calculate the second difference of the time series
## df_stationary = df.diff().diff().dropna()

## Run ADF test on the differenced time series
## result = adfuller(df_stationary['Litecoin'],2)

## Plot the differenced time series
## fig, ax = plt.subplots()
## df_stationary.plot(ax=ax)
## plt.legend(fontsize=8)
## plt.show()

## Print the test statistic and the p-value
## print('ADF Statistic:', result[0])

## Print p-value
## print('p-value:', result[1])

## Print critical values
## print('critical values', result[4])

## Print the test statistic and the p-value
## print('ADF Statistic:', result[0])
```

An alternate method popular for stock price data is log transform.

In [41]: *# Calculate the first difference and drop the nans*

```
df_diff = df.diff().dropna()
```

Run test and print

```
result_diff = adfuller(df['Litecoin'],2)
```

```
print(result_diff)
```

Calculate log-return and drop nans

```
df_log = np.log((df/df.shift(1)).dropna())
```

Run test and print

```
result_log = adfuller(df_log['Litecoin'],2)
```

```
print('p-value:', result_log[1])
```

```
print(result_log)
```

```
fig, ax = plt.subplots()
```

```
df_log.plot(ax=ax)
```

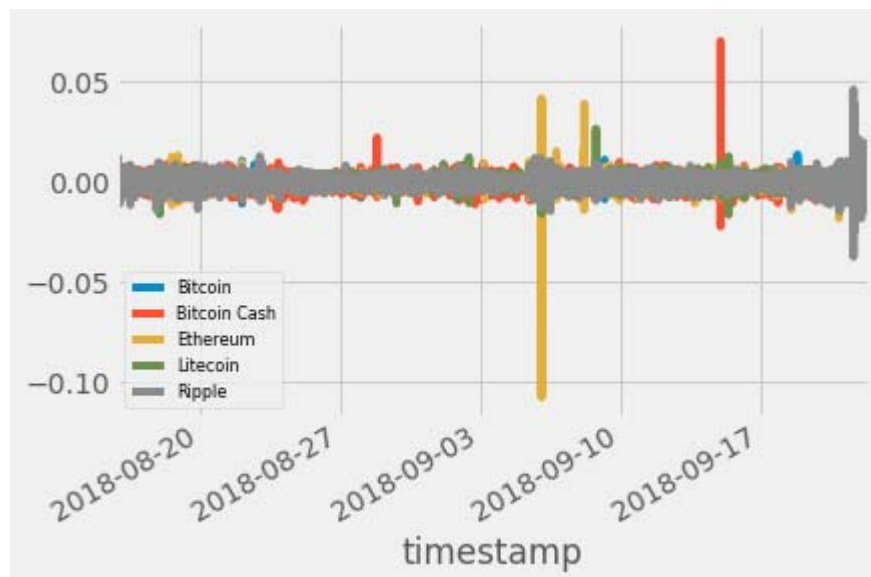
```
plt.legend(fontsize=8)
```

```
plt.show()
```

```
(-2.3526970834847893, 0.1555345950682367, 2L, 657390L, {'5%': -2.861544396638  
8914, '1%': -3.4303599474064623, '10%': -2.5667723401698734}, -3022122.900490  
9424)
```

```
('p-value:', 0.0)
```

```
(-495.578307787468, 0.0, 2L, 657389L, {'5%': -2.8615443966455794, '1%': -3.43  
03599474215942, '10%': -2.566772340173433}, -8341464.317248933)
```



```
In [42]: from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.arima_model import ARMA

# ARIMA model
arima = SARIMAX(df['Litecoin'], order=(1,0,0))

# Fit ARIMA model
arima_results = arima.fit()

# ARIMA forecast of next 10 values
arima_value_forecast = arima_results.get_forecast(steps=10).predicted_mean

# Print forecast
print(arima_value_forecast)

# # Generate predictions
# one_step_forecast = df['Ripple'].get_prediction(start=-30, dynamic=True)

# # Extract prediction mean
# mean_forecast = one_step_forecast.predicted_mean

# # Get confidence intervals of predictions
# confidence_intervals = one_step_forecast.conf_int()

# # Select lower and upper confidence limits
# lower_limits = confidence_intervals.loc[:, 'lower close']
# upper_limits = confidence_intervals.loc[:, 'upper close']

# # Print best estimate predictions
# print(mean_forecast)
```

C:\ProgramData\Anaconda2\lib\site-packages\statsmodels\tsa\base\tsa_model.py:
219: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

' ignored when e.g. forecasting.', ValueWarning)

C:\ProgramData\Anaconda2\lib\site-packages\statsmodels\tsa\base\tsa_model.py:
576: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

ValueWarning)

```
657393    58.759996
657394    58.759992
657395    58.759988
657396    58.759983
657397    58.759979
657398    58.759975
657399    58.759971
657400    58.759967
657401    58.759963
657402    58.759959
dtype: float64
```

In []: